

How to build an Event-Sourcing system using Akka with EKS

ScalaMatsuri 2019

Junichi Kato(@j5ik2o)



Who am I

- Chatwork Tech-Lead
- github/j5ik2o
 - [scala-ddd-base](#)
 - [scala-ddd-base-akka-http.g8](#)
 - [reactive-redis](#)
 - [reactive-memcached](#)
- Review
 - [Domain-Driven Design: Tackling Complexity in the Heart of Software \(Japanese Edition\)](#)
 - [Akka in Action \(Japanese Edition\)](#)



Agenda

1. Event Sourcing with Akka
 2. Deployment to EKS
- <https://github.com/j5ik2o/thread-weaver>

Akka with Event Sourcing

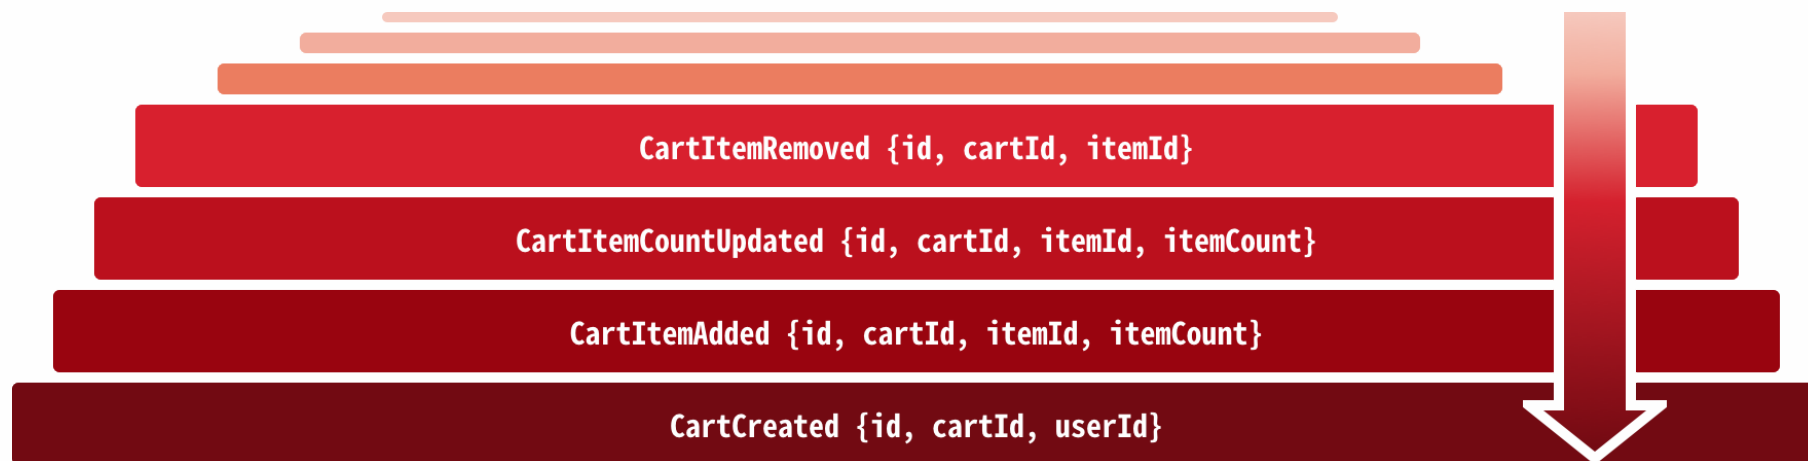
Event Sourcing

- The latest state is derived by the events
- For example, transactions such as the e-commerce are sourced on events. This is nothing special.
- An event sequence represents an immutable history.
 - The transaction makes the following unique corrections. Events are never modified or deleted.
 - The order #0001 is canceled at the #0700, and the corrected data is registered at the slip #0701.

Slip Number	Product	Price	Quantity	Slip Number For Correction	Remarks
0001	A0123	5,000	10	0700	data before modification
0700	A0123	5,000	-10	0001	data for cancellation
0701	A0123	4,000	20		corrected data

Domain Events

- Events that occurred in the past
- Domain Events are events that interest domain experts.
- Generally, Domain Events is expressed as a verb in past tense
 - CustomerRelocated
 - CargoShipped
- Events and commands are similar, but different languages are handled by humans
 - Command may be rejected
 - Indicates that the event has already occurred

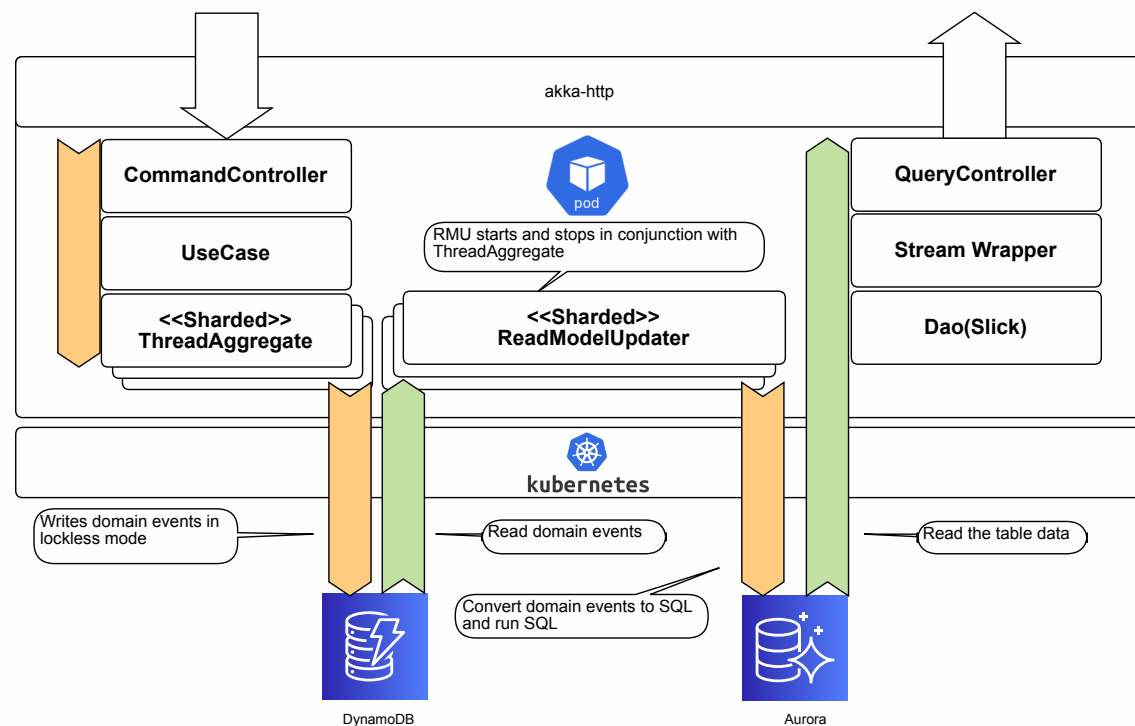


**Consider thread-weaver
as an example of a simple chat application.**

System requirements

- API server accepts commands and queries from API clients
- Create a thread to start the chat
- Only members can post to threads
- Only text messages posted to threads
- Omit authentication and authorization for convenience

System Configuration

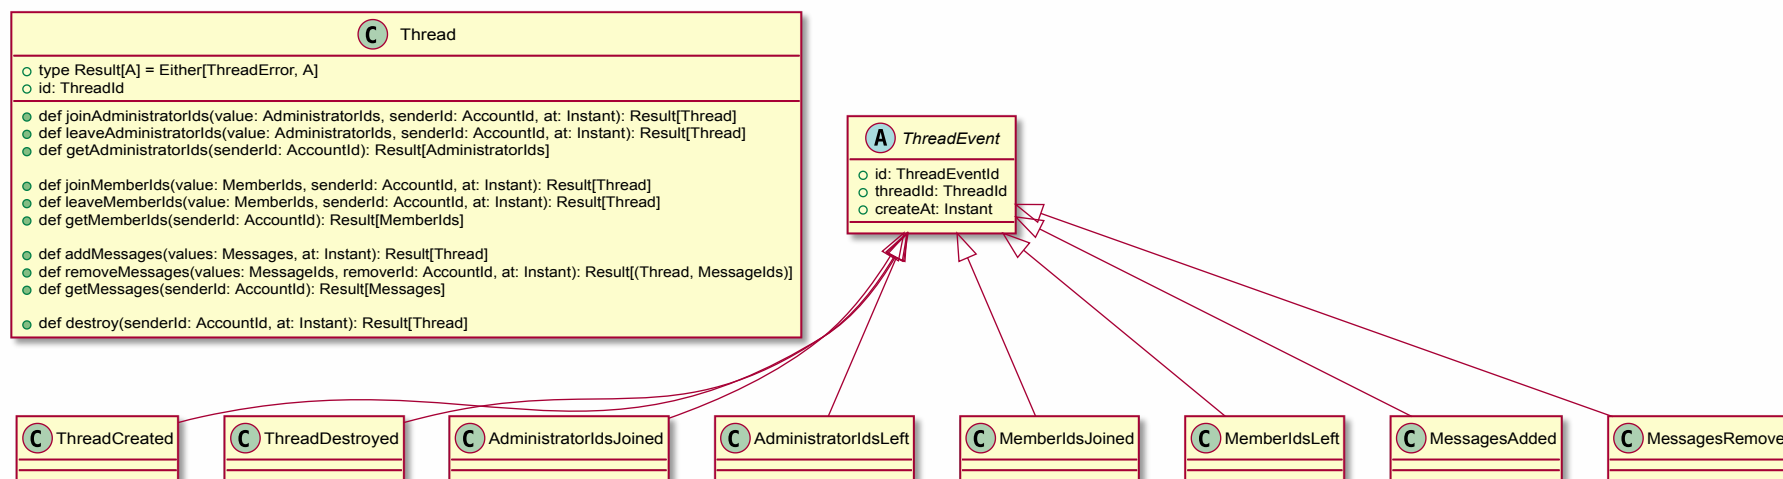


- Split the application into the command stack and the query stack
- The command is sent to (clustered sharding) aggregate actor
- The aggregate actor stores(appends) domain events in storage when it accepts a command
- RMU(cluster sharding) starts up in conjunction with the aggregation actor and reads the domain events for the appropriate aggregate ID immediately after startup, executes the SQL, and creates the Read-Model
- Query using DAO to load and return the lead model
- Deploy the api-server as a kubernetes pod

Command stack side

Domain Objects

- Account
 - Account information identifying the user of the system
- Thread
 - Indicates a place to exchange Messages
- Message
 - A hearsay written in some language
- Administrator
 - Administrator of the Thread
- Member
 - Users of the Thread



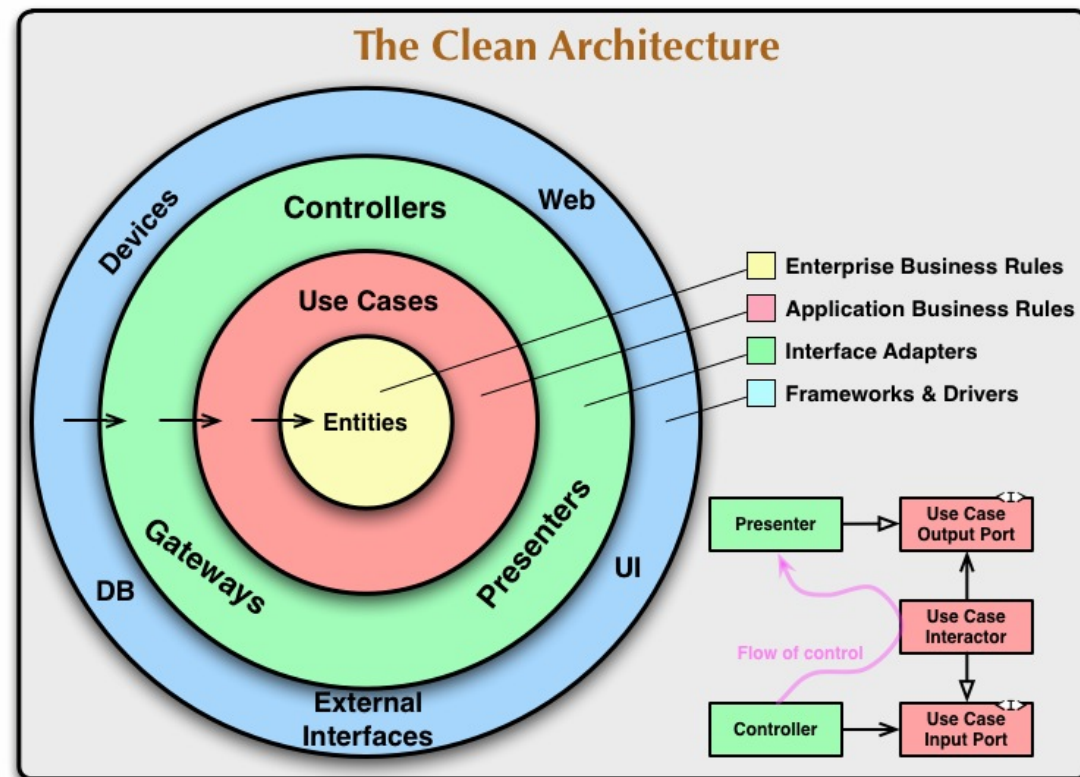
Commands/Domain Events

ThreadEvent sub types

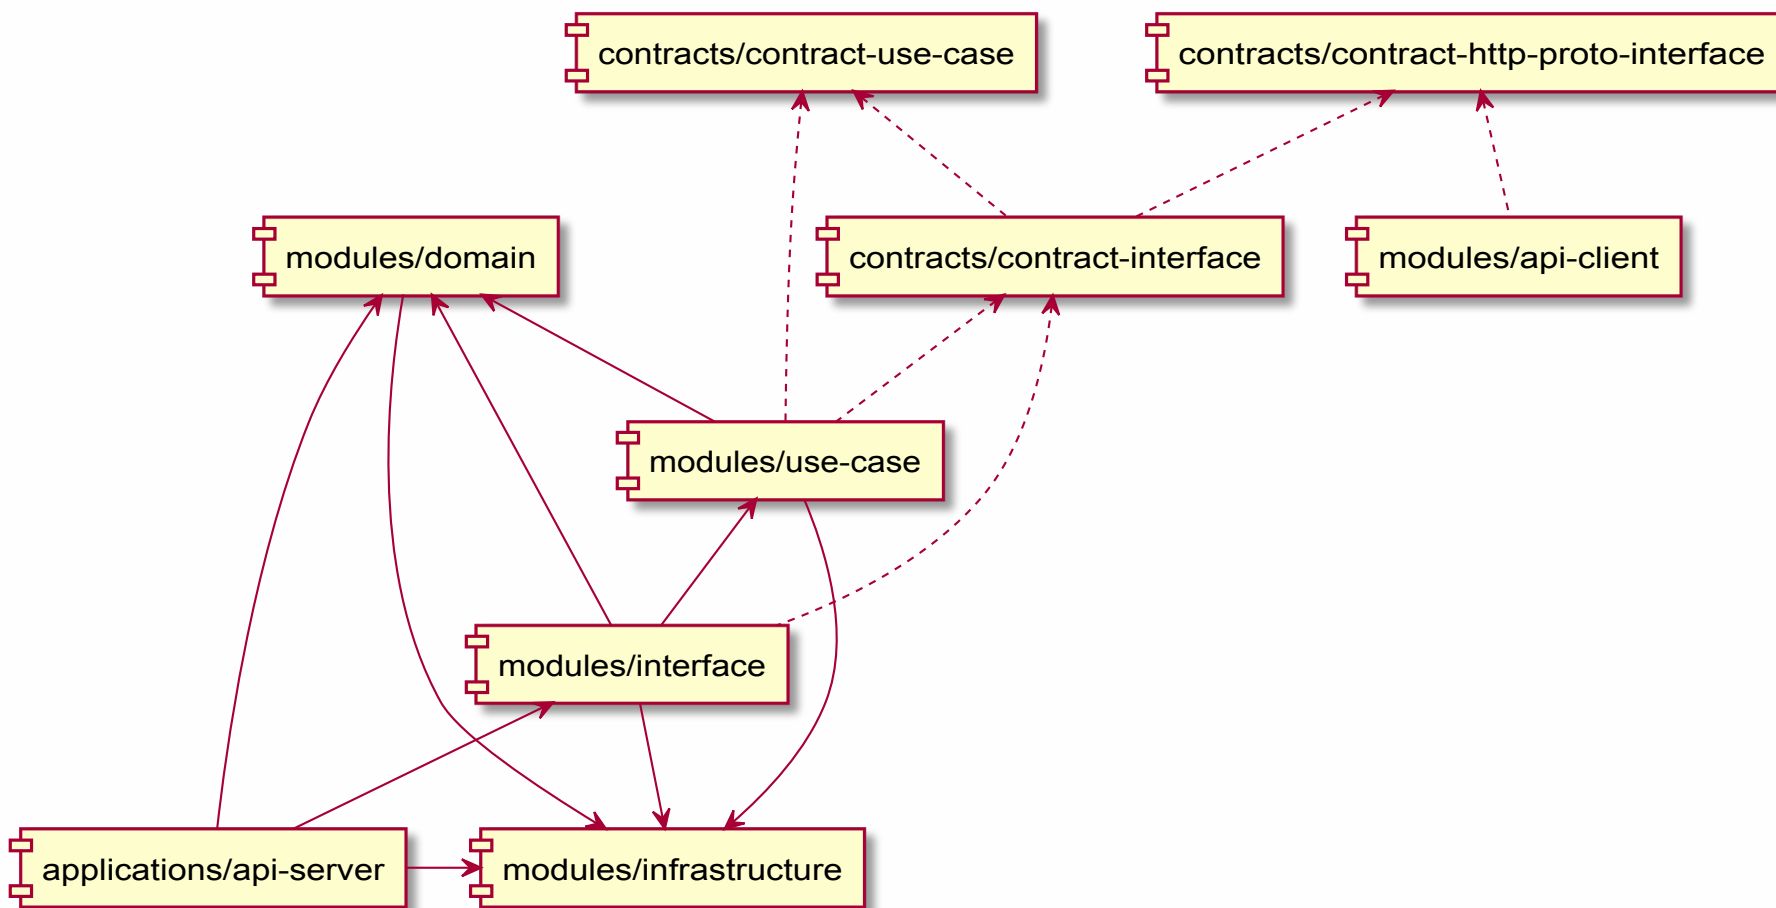
- Create/Destroy Thread
 - ThreadCreated
 - ThreadDestroyed
- Join/Leave AdministratorIds
 - AdministratorIdsJoined
 - AdministratorIdsLeft
- Join/Leave MemberIds
 - MemberIdsJoined
 - MemberIdsLeft
- Add/Remove Messages
 - MessagesAdded
 - MessagesRemoved

Layered architecture

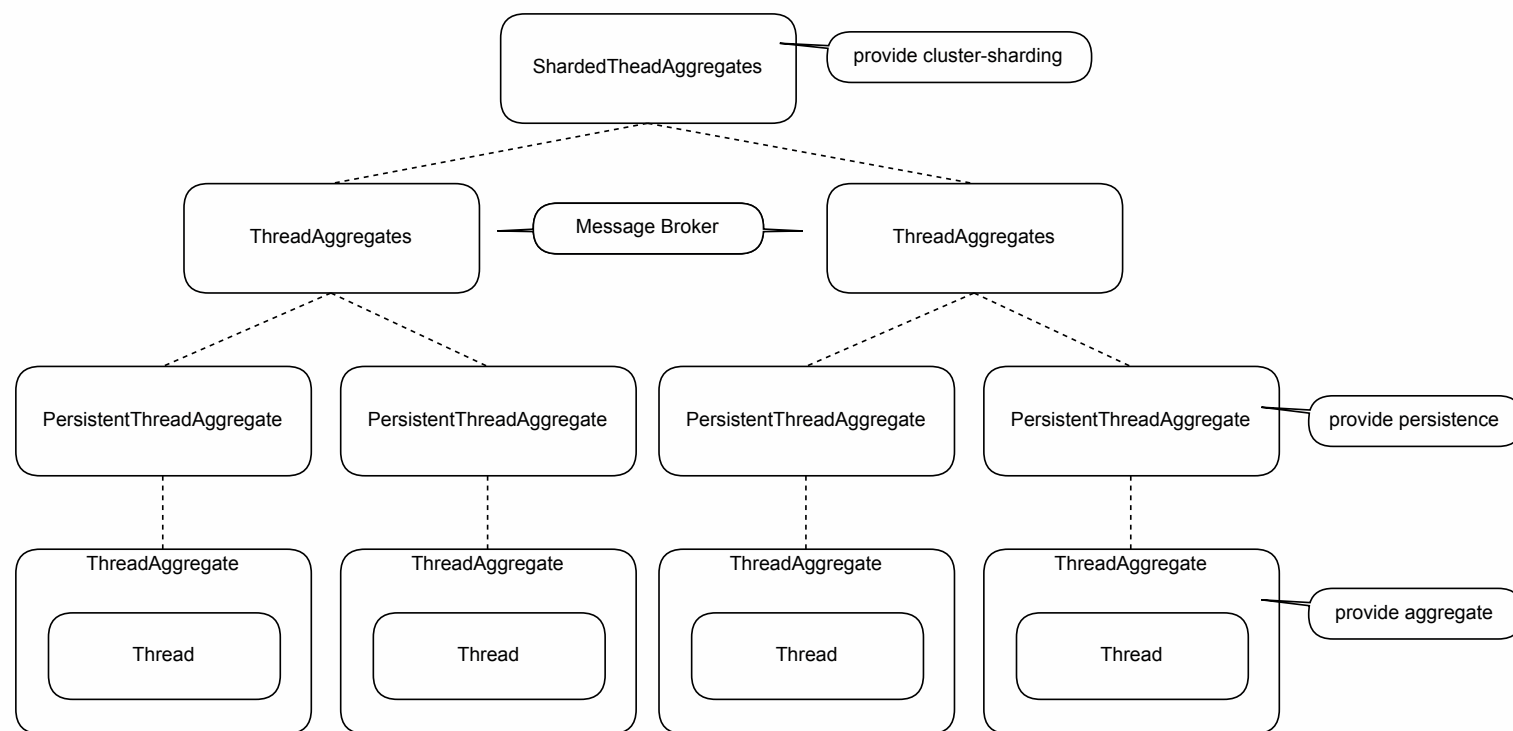
- Clean Architecture
 - interface-adaptors
 - infrastructure
- Command side
 - use-cases
 - domain
- Query side
 - data access streams
 - data access objects



Projects structure



Domain objects with actors



- Actors that fulfill all the functions are undesirable
- Follow object-oriented principles to build a hierarchy of actors with a single responsibility

Thread

```
trait Thread {  
  
  def isAdministratorId(accountId: AccountId): Boolean  
  def isMemberId(accountId: AccountId): Boolean  
  
  def joinAdministratorIds(value: AdministratorIds, senderId: AccountId, at: Instant): Result[Thread]  
  def leaveAdministratorIds(value: AdministratorIds, senderId: AccountId, at: Instant): Result[Thread]  
  def getAdministratorIds(senderId: AccountId): Result[AdministratorIds]  
  
  def joinMemberIds(value: MemberIds, senderId: AccountId, at: Instant): Result[Thread]  
  def leaveMemberIds(value: MemberIds, senderId: AccountId, at: Instant): Result[Thread]  
  def getMemberIds(senderId: AccountId): Result[MemberIds]  
  
  def addMessages(values: Messages, at: Instant): Result[Thread]  
  def removeMessages(values: MessageIds, removerId: AccountId, at: Instant): Result[(Thread, MessageIds)]  
  def getMessages(senderId: AccountId): Result[Messages]  
  
  def destroy(senderId: AccountId, at: Instant): Result[Thread]  
}
```


ThreadAggregate

```
class ThreadAggregate(id: ThreadId,
  subscribers: Seq[ActorRef]) extends Actor {
  // add messages handler
  private def commandAddMessages(thread: Thread): Receive = {
    case AddMessages(requestId, threadId,
      messages, createdAt, reply) if threadId == id =>
      thread.addMessages(messages, createdAt) match {
        case Left(exception) =>
          if (reply)
            sender() ! AddMessagesFailed(ULID(), requestId,
              threadId, exception.getMessage, createdAt)
        case Right(newThread) =>
          if (reply)
            sender() ! AddMessagesSucceeded(ULID(), requestId,
              threadId, messages.toMessageIds, createdAt)
          context.become(onCreated(newThread))
      }
  }

  override def receive: Receive = { /*...*/ }
}
```

- Actors that support transactional integrity
- The boundary of the data update is the same as the boundary the aggregates has
- For example, when an actor receives CreateThread command, a Thread state is generated internally
- Then Messages are also added to the Thread when the AddMessages command is received
- If the other commands defined in the protocol are received by the Actor, the Actor will have corresponding side effects.

ThreadAggregateSpec

```
val threadId      = ThreadId()
val threadRef     = newThreadRef(threadId)
val now          = Instant.now
val administratorId = AccountId()
val title        = ThreadTitle("test")

threadRef ! CreateThread(ULID(), threadId, administratorId,
  None, title, None, AdministratorIds(administratorId),
  MemberIds.empty, now, reply = false)

val messages = Messages(TextMessage(MessageId(), None,
  ToAccountIds.empty, Text("ABC"), memberId, now, now))
threadRef ! AddMessages(ULID(), threadId, messages,
  now, reply = true)

expectMsgType[AddMessagesResponse] match {
  case f: AddMessagesFailed =>
    fail(f.message)
  case s: AddMessagesSucceeded =>
    s.threadId shouldBe threadId
    s.createAt shouldBe now
}
```

- Verify that add messages and create a thread by using Test Kit

PersistentThreadAggregate(1/2)

```
object PersistentThreadAggregate {
  def props(id: ThreadId, subscribers: Seq[ActorRef]): Props =
    ...
}

class PersistentThreadAggregate(id: ThreadId,
  subscribers: Seq[ActorRef],
  propsF: (ThreadId, Seq[ActorRef]) => Props)
  extends PersistentActor with ActorLogging {

  override def supervisorStrategy: SupervisorStrategy =
    OneForOneStrategy() { case _: Throwable => Stop }

  private val childRef =
    context.actorOf(propsF(id, subscribers),
      name = ThreadAggregate.name(id))
    context.watch(childRef)

  // persistenceId is the partition keyでもスレッドIDに対応するf persi

  override def receiveRecover: Receive = {
    case e: ThreadCommonProtocol.Event with ToCommandRequest =>
      childRef ! e.toCommandRequest
    case RecoveryCompleted =>
      log.debug("recovery completed")
  }
}
```

- Actors that add the persistence function to ThreadAggregate
- Domain behavior is provided by child actors
- The recover process sends commands generated from events to child actors.

PersistentThreadAggregate(2/2)

```
override def receiveCommand: Receive = {
  case Terminated(c) if c == childRef =>
    context.stop(self)
  case m: CommandRequest with ToEvent =>
    childRef ! m
    context.become(sending(sender(), m.toEvent))
  case m =>
    childRef forward m
}

private def sending(replyTo: ActorRef,
  event: ThreadCommonProtocol.Event): Receive = {
  case s: CommandSuccessResponse => persist(event) { _ =>
    replyTo ! s
    unstashAll()
    context.unbecome()
  }
  case f: CommandFailureResponse =>
    replyTo ! f
    unstashAll()
    context.unbecome()
  case _ =>
    stash()
}
```

- Delegate to child actors when receiving commands. Persists only on success
- message processing is suspended until a command response is returned

PersitentThreadAggregateSpec

```
// Create id = 1 of Thread actor
threadRef1 ! CreateThread(ULID(), threadId, administratorId, None, title, None,
  AdministratorIds(administratorId), MemberIds.empty, now, reply = false)
val messages = Messages(TextMessage(MessageId(), None,
  ToAccountIds.empty, Text("ABC"), memberId, now, now))
threadRef1 ! AddMessages(ULID(), threadId, messages, now, reply = false)

//Stop id = 1 of Thread actor
killActors(threadRef)

// Recover id = 1 of Thread actor
val threadRef2 = system.actorOf(PersistentThreadAggregate.props(threadId, Seq.empty))

// Check if it is in the previous state
threadRef2 ! GetMessages(ULID(), threadId, memberId, now)
expectMsgType[GetMessagesResponse] match {
  case f: GetMessagesFailed =>
    fail(f.message)
  case s: GetMessagesSucceeded =>
    s.threadId shouldBe threadId
    s.createAt shouldBe now
    s.messages shouldBe messages
}
```

- a test that intentionally stops and restarts the persistence actor
- Replayed state after reboot

FYI: akka-persistence plugin

- Different plugins for different databases
 - [Akka Persistence journal and snapshot plugins](#)
- Default corresponds to LevelDB
- recommended on AWS is DynamoDB. There are the following plugins, but I recommend my plugin:P
 - <https://github.com/j5ik2o/akka-persistence-dynamodb>
 - <https://github.com/akka/akka-persistence-dynamodb>

FYI: akka-persistence plugin

- akka plugin
 - is complex schema
 - doesn't support the query function.
 - doesn't support non-blocking I/O (by aws-sdk v2)

akka/akka-persistence-dynamodb

```
# PersistentRepr
PKey: par = <journalName>-P-<persistenceId>-<sequenceNr / 100>
SKey: num = <sequenceNr % 100>
pay = <payload>
idx = <atomic write batch index>
cnt = <atomic write batch max index>
# High sequence number
PKey: par = <journalName>-SH-<persistenceId>-<(sequenceNr / 100)
    % sequenceShards>
SKey: num = 0
seq = <sequenceNr rounded down to nearest multiple of 100>
# Low sequence number
PKey: par = <journalName>-SL-<persistenceId>-<(sequenceNr / 100)
    % sequenceShards>
SKey: num = 0
seq = <sequenceNr, not rounded>
```

j5ik2o/akka-persistence-dynamodb

```
# PrimaryIndex(for Writing)
PKey: pkey = ${PersistenceId}-${SequenceNumber % ShardCount}
SKey: sequence-nr = ${SequenceNumber}
# GSI:GetJournalRows(for Reading)
persistence-id = ${PersistenceId}
sequence-nr = ${SequenceNumber}
# GSI:TagsIndex(for Reading)
tags = ...
```

FYI: akka-persistence plugin

```
// build.sbt
libraryDependencies += Seq(
  // ...
  "com.typesafe.akka" %% "akka-persistence" % akkaVersion,
  "org.fusesource.leveldbjni" % "leveldbjni-all" % "1.8" % Test,
  "org.iq80.leveldb" % "leveldb" % "0.9" % Test,
  "com.github.j5ik2o" %% "akka-persistence-dynamodb" % "1.0.2",
  // ...
)
```

```
// application.conf
akka {
  persistence {
    journal {
      plugin = dynamo-db-journal
    }
    snapshot-store {
      plugin = dynamo-db-snapshot
    }
  }
}
```


ThreadAggregates(Message Broker)

```
class ThreadAggregates(subscribers: Seq[ActorRef],
  propsF: (ThreadId, Seq[ActorRef]) => Props)
  extends Actor
  with ActorLogging
  with ChildActorLookup {
  override type ID = ThreadId
  override type CommandRequest = ThreadProtocol.CommandMessage

  override def receive: Receive = forwardToActor

  override protected def childName(childId: ThreadId): String =
    childId.value.asString
  override protected def childProps(childId: ThreadId): Props =
    propsF(childId, subscribers)
  override protected def toChildId(commandRequest: CommandMessage): ThreadId =
    commandRequest.threadId
}
```

- The message broker that bundles multiple ThreadAggregates as child actors
- Most of the logic is in ChildActorLookup
- Resolve the actor name from ThreadId in the command message, and transfer the message to the corresponding child actor. If there is no child actor, generate an actor and then forward the message to the actor

ChildActorLookup

```
trait ChildActorLookup extends ActorLogging { this: Actor =>
  implicit def context: ActorContext
  type ID
  type CommandRequest

  protected def childName(childId: ID): String
  protected def childProps(childId: ID): Props
  protected def toChildId(commandRequest: CommandRequest): ID

  protected def forwardToActor: Actor.Receive = {
    case _cmd =>
      val cmd = _cmd.asInstanceOf[CommandRequest]
      context
        .child(childName(toChildId(cmd)))
        .fold(createAndForward(cmd, toChildId(cmd)))(forwardCommand(cmd))
  }

  protected def forwardCommand(cmd: CommandRequest)(childRef: ActorRef): Unit =
    childRef forward cmd

  protected def createAndForward(cmd: CommandRequest, childId: ID): Unit =
    createActor(childId) forward cmd

  protected def createActor(childId: ID): ActorRef =
    context.actorOf(childProps(childId), childName(childId))
}
```

- Create a child actor if none exists and forward the message
- forward the message to its child actors, if any

ShardedThreadAggregates (1/2)

```
object ShardedThreadAggregates {  
  def props(subscribers: Seq[ActorRef],  
    propsF: (ThreadId, Seq[ActorRef]) => Props): Props =  
    Props(new ShardedThreadAggregates(subscribers, propsF))  
  
  def name(id: ThreadId): String = id.value.asString  
  
  val shardName = "threads"  
  
  case object StopThread  
  
  // function to extract an entity id  
  val extractEntityId: ShardRegion.ExtractEntityId = {  
    case cmd: CommandRequest => (cmd.threadId.value.asString, cmd)  
  }  
  
  // function to extract a shard id  
  val extractShardId: ShardRegion.ExtractShardId = {  
    case cmd: CommandRequest =>  
      val mostSignificantBits = cmd.threadId  
        .value.mostSignificantBits % 12  
      val leastSignificantBits = cmd.threadId  
        .value.leastSignificantBits % 12  
      s"$mostSignificantBits:$leastSignificantBits"  
  }  
}
```

- Allow ThreadAggregates to be distributed across a cluster
- extractEntityId is the function to extract an entity id
- extractShardId is the function to extract a shard id

ShardedThreadAggregates (2/2)

```
class ShardedThreadAggregates(subscribers: Seq[ActorRef],  
  propsF: (ThreadId, Seq[ActorRef]) => Props)  
  extends ThreadAggregates(subscribers, propsF) {  
    context.setReceiveTimeout(  
      Settings(context.system).passivateTimeout)  
  
    override def unhandled(message: Any): Unit = message match {  
      case ReceiveTimeout =>  
        log.debug("ReceiveTimeout")  
        context.parent ! Passivate(stopMessage = StopThread)  
      case StopThread =>  
        log.debug("StopWallet")  
        context.stop(self)  
    }  
  }
```

- Inherit ThreadAggregates
- Then add an implementation to passivate ShardedThreadAggregates when occurred ReceiveTimeout

ShardedThreadAggregatesRegion

```
object ShardedThreadAggregatesRegion {  
  def startClusterSharding(subscribers: Seq[ActorRef])  
    (implicit system: ActorSystem): ActorRef =  
    ClusterSharding(system).start(  
      ShardedThreadAggregates.shardName,  
      ShardedThreadAggregates.props(subscribers,  
        PersistentThreadAggregate.props),  
      ClusterShardingSettings(system),  
      ShardedThreadAggregates.extractEntityId,  
      ShardedThreadAggregates.extractShardId  
    )  
  
  def shardRegion(implicit system: ActorSystem): ActorRef =  
    ClusterSharding(system)  
      .shardRegion(ShardedThreadAggregates.shardName)  
}
```

- The startClusterSharding method will start ClusterSharing with the specified settings
- The shardRegion method gets the ActorRef to the started ShardRegion.

MultiJVM Testing(by using sbt-multi-jvm)

```

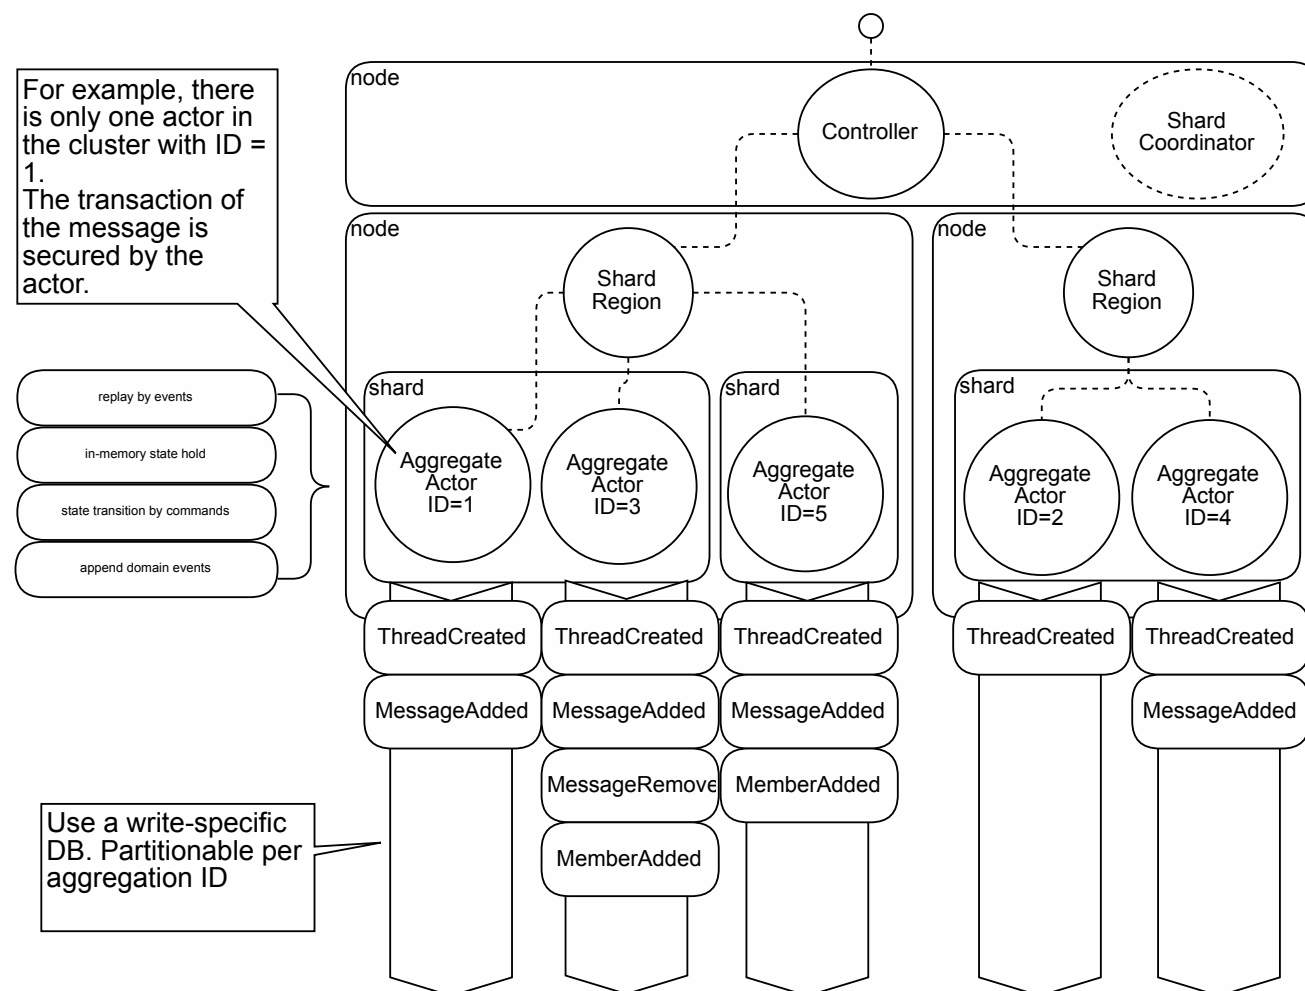
"setup shared journal" in {
  Persistence(system)
  runOn(controller) { system.actorOf(Props[SharedLevelDbStore], "store") }
  enterBarrier("persistence-started")
  runOn(node1, node2) {
    system.actorSelection(node(controller) / "user" / "store") ! Identify(None)
    val sharedStore = expectMsgType[ActorIdentity].ref.get
    SharedLevelDbJournal.setStore(sharedStore, system)
  }
  enterBarrier("setup shared journal")
}

"join cluster" in within(15 seconds) {
  join(node1, node1) { ShardedThreadAggregatesRegion.startClusterSharding(Seq.empty) }
  join(node2, node1) { ShardedThreadAggregatesRegion.startClusterSharding(Seq.empty) }
  enterBarrier("join cluster")
}

"createThread" in { runOn(node1) {
  val accountId = AccountId(); val threadId = ThreadId(); val title = ThreadTitle("test")
  val threadRef = ShardedThreadAggregatesRegion.shardRegion
  threadRef ! CreateThread(ULID(), threadId, accountId, None, title, None, AdministratorIds(accountId),
    MemberIds.empty, Instant.now, reply = true)
  expectMsgType[CreateThreadSucceeded](file:///Users/j5ik2o/Sources/thread-weaver/slide/10 seconds).threadId shouldBe threadId
}
  enterBarrier("create thread")
}

```

cluster-sharding with persistence



- Actors with state in on-memory are distributed across the cluster
- Domain events that occur are saved in partitioned storage by aggregate ID

CreateThreadUseCaseUntypedImpl

```
class CreateThreadUseCaseUntypedImpl(
  threadAggregates: ThreadActorRefOfCommandUntypedRef, parallelism: Int = 1, timeout: Timeout = 3 seconds
)(implicit system: ActorSystem) extends CreateThreadUseCase {
  override def execute: Flow[UCreateThread, UCreateThreadResponse, NotUsed] =
    Flow[UCreateThread].mapAsync(parallelism) { request =>
      implicit val to: Timeout = timeout
      implicit val scheduler: Scheduler = system.scheduler
      implicit val ec: ExecutionContextExecutor = system.dispatcher
      (threadAggregates ? CreateThread(
        ULID(), request.threadId, request.creatorId, None, request.title, request.remarks,
        request.administratorIds, request.memberIds, request.createAt, reply = true
      )).mapTo[CreateThreadResponse].map {
        case s: CreateThreadSucceeded =>
          UCreateThreadSucceeded(s.id, s.requestId, s.threadId, s.createAt)
        case f: CreateThreadFailed =>
          UCreateThreadFailed(f.id, f.requestId, f.threadId, f.message, f.createAt)
      }
    }
}
```


ThreadCommandControllerImpl

```
trait ThreadCommandControllerImpl
  extends ThreadCommandController
  with ThreadValidateDirectives {
  private val createThreadUseCase = bind[CreateThreadUseCase]
  private val createThreadPresenter = bind[CreateThreadPresenter]

  override private[controller] def createThread: Route =
    path("threads" / "create") {
      post {
        extractMaterializer { implicit mat =>
          entity(as[CreateThreadRequestJson]) { json =>
            validateJsonRequest(json).apply { commandRequest =>
              val responseFuture = Source.single(commandRequest)
                .via(createThreadUseCase.execute)
                .via(createThreadPresenter.response)
                .runWith(Sink.head)
              onSuccess(responseFuture) { response =>
                complete(response)
              }
            }
          }
        }
      }
    }
  }
}
```

- Command side controller
- The thread creation root composes several directives and calls a use case
- The request JSON returns a command if validation passes. Pass the command to the use-case and execute it
- The presenter will convert the use-case result to Response JSON

Read Model Updater side

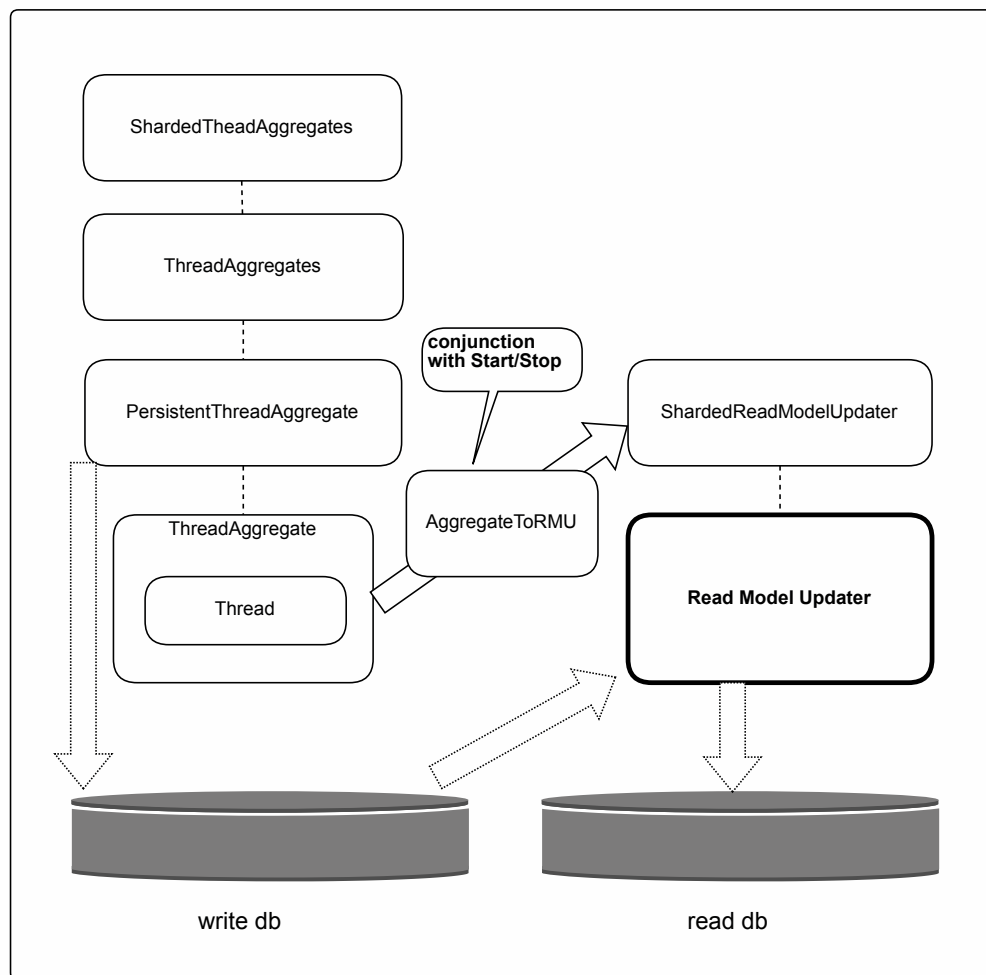
FYI: akka-typed

- The message type received by the message handler was Any, but akka-typed allows the Message type to be specified
- There is basically no compatibility, so there are many things to remember. Let's get used to it now

```
object PingPong extends App {  
  trait Message  
  case class Ping(reply: ActorRef[Message]) extends Message  
  case object Pong extends Message  
  
  def receiver: Behavior[Message] =  
    Behaviors.setup[Message] { ctx =>  
      Behaviors.receiveMessagePartial[Message] {  
        case Ping(replyTo) =>  
          ctx.log.info("ping")  
          replyTo ! Pong  
          Behaviors.same  
      }  
    }  
}
```

```
def main: Behavior[Message] = Behaviors.setup { ctx =>  
  val receiverRef = ctx.spawn(receiver, name = "receiver")  
  receiverRef ! Ping(ctx.self)  
  Behaviors.receiveMessagePartial[Message] {  
    case Pong =>  
      ctx.log.info("pong")  
      receiverRef ! Ping(ctx.self)  
      Behaviors.same  
  }  
}  
  
ActorSystem(main, "ping-pong")  
}
```

Read Model Updater(1/3)



- Starts the Read Model Updater (RMU) for each aggregation ID
- Sharding to allow multiple RMUs to boot on a single node
- Starting and stopping the RMU is triggered by events on the aggregate actor. It actually does message translation with `AggregateToRMU`.

Read Model Updater(2/3)

```
private def projectionSource(sqlBatchSize: Long, threadId: ThreadId)
  (implicit ec: ExecutionContext): Source[Vector[Unit], NotUsed] = {
  Source
    .fromFuture(
      db.run(getSequenceNrAction(threadId)).map(_._getOrElse(0L))
    ).log("lastSequenceNr").flatMapConcat { lastSequenceNr =>
      readJournal
        .eventsByPersistenceId(threadId.value.asString,
          lastSequenceNr + 1, Long.MaxValue)
    }.log("ee").via(sqlActionFlow(threadId))
    .batch(sqlBatchSize, ArrayBuffer(_))(_ :+ _).mapAsync(1) { sqlActions =>
      db.run(DBIO.sequence(sqlActions.result.toVector))
    }.withAttributes(logLevels)
}
```

- RMU does not end stream processing - persistenceId also gets the latest sequence number corresponding to the thread ID
- read events from readJournal since thread ID and last sequence number
- sqlActionFlow converts events to SQL
- Finally, run the SQL in batches (Read model not denormalized to be flexible to query patterns)

Read Model Updater(2/2)

```
class ThreadReadModelUpdater(
  val readJournal: ReadJournalType,
  val profile: JdbcProfile, val db: JdbcProfile#Backend#Database
) extends ThreadComponent with ThreadMessageComponent ... {
  import profile.api._

  def behavior(sqlBatchSize: Long = 10,
    backoffSettings: Option[BackoffSettings] = None): Behavior[CommandRequest] =
    Behaviors.setup[CommandRequest] { ctx =>
      Behaviors.receiveMessagePartial[CommandRequest] {
        case s: Start =>
          ctx.child(s.threadId.value.asString) match {
            case None =>
              ctx.spawn(
                projectionBehavior(sqlBatchSize, backoffSettings, s.threadId),
                name = s"RMU-${s.threadId.value.asString}"
              ) ! s
            case _ =>
              ctx.log.warning(
                "RMU already has started: threadId = {}", s.threadId.value.asString)
          }
        Behaviors.same
      }
    }
  // ...
}
```

- RMU starts stream processing when it receives a Start message.
- Stream processing is performed as a task on a child actor

ShardedThreadReadModelUpdater(1/3)

```
class ShardedThreadReadModelUpdater(  
  val readJournal: ReadJournalType,  
  val profile: JdbcProfile,  
  val db: JdbcProfile#Backend#Database  
) {  
  val TypeKey: EntityTypeKey[CommandRequest] = EntityTypeKey[CommandRequest](file:///Users/j5ik2o/Sources/thread-weaver/slide/"threads  
  
  def initEntityActor(  
    clusterSharding: ClusterSharding,  
    receiveTimeout: FiniteDuration  
  ): ActorRef[ShardingEnvelope[CommandRequest]] =  
    clusterSharding.init(  
      Entity(typeKey = TypeKey, createBehavior = behavior(receiveTimeout)).withStopMessage(Stop)  
    )  
  // ...
```

ShardedThreadReadModelUpdater(2/3)

```
// ...
private def behavior(
  receiveTimeout: FiniteDuration, sqlBatchSize: Long = 10, backoffSettings: Option[BackoffSettings] = None
): EntityContext => Behavior[CommandRequest] = { entityContext =>
  Behaviors.setup[CommandRequest] { ctx =>
    // setting receive timeout
    ctx.setReceiveTimeout(receiveTimeout, Idle)
    val childRef = ctx.spawn(new ThreadReadModelUpdater(readJournal, profile, db).behavior(sqlBatchSize, backoffSettings),
      name = "threads-rmu")
    Behaviors.receiveMessagePartial {
      case Idle => entityContext.shard ! ClusterSharding.Passivate(ctx.self); Behaviors.same
      case Stop => Behaviors.stopped
      case msg => childRef ! msg; Behaviors.same
    }
  }
}
// ...
```


ShardedThreadReadModelUpdater(3/3)

- MUST send the ShardingEnvelope to the ShardRegion. ShadedThreadReadModelUpdaterProxy converts CommandRequest into ShardingEnvelope and forwards it

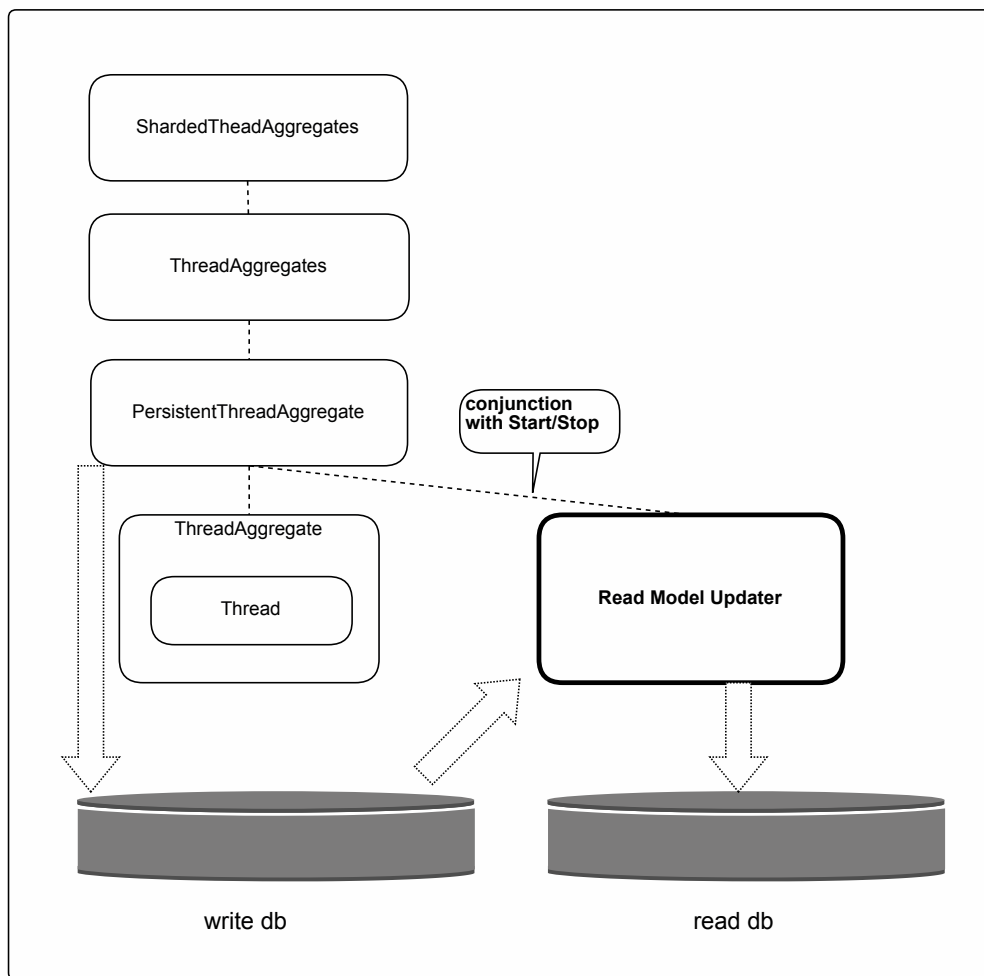
```
class ShardedThreadReadModelUpdaterProxy(  
  val readJournal: ReadJournalType, val profile: JdbcProfile, val db: JdbcProfile#Backend#Database  
) {  
  def behavior(clusterSharding: ClusterSharding, receiveTimeout: FiniteDuration): Behavior[CommandRequest] =  
    Behaviors.setup[CommandRequest] { _ =>  
      val actorRef =  
        new ShardedThreadReadModelUpdater(readJournal, profile, db).initEntityActor(clusterSharding, receiveTimeout)  
      Behaviors.receiveMessagePartial[CommandRequest] {  
        case msg =>  
          actorRef ! typed.ShardingEnvelope(msg.threadId.value.asString, msg)  
          Behaviors.same  
        }  
      }  
    }  
}
```

Interlocking of Aggregate and RMU

```
object AggregateToRMU {
  def behavior(
    rmuRef: ActorRef[ThreadReadModelUpdaterProtocol.CommandRequest]
  ): Behavior[ThreadCommonProtocol.Message] =
    Behaviors.setup[ThreadCommonProtocol.Message] { ctx =>
      Behaviors.receiveMessagePartial[ThreadCommonProtocol.Message] {
        case s: Started =>
          ctx.log.debug(s"RMU ! $s")
          rmuRef ! ThreadReadModelUpdaterProtocol.Start(
            ULID(), s.threadId, Instant.now)
          Behaviors.same
        case s: Stopped =>
          ctx.log.debug(s"RMU ! $s")
          rmuRef ! ThreadReadModelUpdaterProtocol.Stop(
            ULID(), s.threadId, Instant.now)
          Behaviors.same
      }
    }
}
```

- These two actors are separated because they have different responsibilities, but start and stop work together
- Actually there is a problem with this method. If only the RMU stops due to a node failure, the RMU cannot recover until it receives the Start message again. The downside is that ThreadAggregate must periodically send heartbeat beads.

Improvement for RMU



- Another implementation pattern is to make the RMU a child actor of PersistentThreadAggregate.
- This method allows you to watch the RMU as a child actor so that it can be restarted if the RMU should stop.
- However, PersistentThreadAggregate is responsible for RMU responsibilities. Duplicate Responsibilities?

Query stack side

ThreadQueryControllerImpl

```

trait ThreadQueryControllerImpl
  extends ThreadQueryController
  with ThreadValidateDirectives {
  private val threadDas: ThreadDas = bind[ThreadDas]
  // ...

  override private[controller] def getThread: Route =
    path("threads" / Segment) { threadIdString => get {
      // ...
      parameter('account_id) { accountValue =>
        validateAccountId(accountValue) { accountId =>
          onSuccess(threadDas.getThreadByIdSource(accountId, threadId)
            .via(threadPresenter.response)
            .runWith(Sink.headOption[ThreadJson]).map(identity)) {
            case None =>
              reject(NotFoundRejection("thread is not found", None))
            case Some(response) =>
              complete(GetThreadResponseJson(response))
          }
        }
      }
    }
  }
  // ...
}

```

- The query side uses a stream wrapped Dao object instead of a use case.
- Same as command side except for this.

ThreadControllerSpec

```

val administratorId = ULID().asString
val entity = CreateThreadRequestJson(
  administratorId, None, "test",
  None, Seq(administratorId), Seq.empty,
  Instant.now.toEpochMilli
).toHttpEntity

Post(RouteNames.CreateThread, entity) ~>
  commandController.createThread ~> check {
    response.status shouldEqual StatusCodes.OK
    val responseJson = responseAs[CreateThreadResponseJson]
    responseJson.isSuccessful shouldBe true
    val threadId = responseJson.threadId.get

    eventually { // repeat util read
      Get(RouteNames.GetThread(threadId, administratorId)) ~>
        queryController.getThread ~> check {
          response.status shouldEqual StatusCodes.OK
          val responseJson = responseAs[GetThreadResponseJson]
          responseJson.isSuccessful shouldBe true
        }
    }
  }
}

```

- a test where two controllers are connected.
- Verify threads are readable after they are created
- Works fine

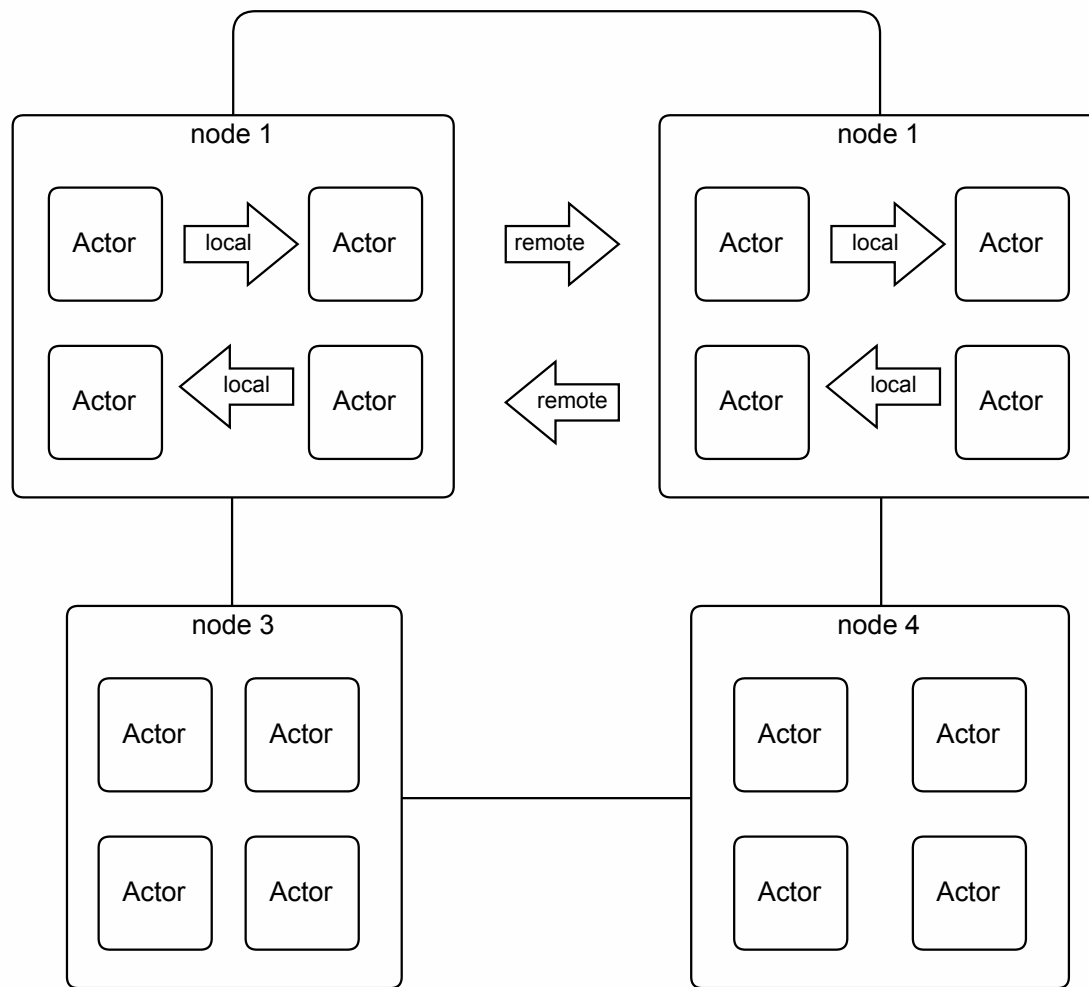
Bootstrap

- Start AkkaManagement and ClusterBootstrap, then start akka-http server

```
object Main extends App {  
  // ...  
  
  implicit val system: ActorSystem                = ActorSystem("thread-weaver-api-server", config)  
  implicit val materializer: ActorMaterializer    = ActorMaterializer()  
  implicit val executionContext: ExecutionContext = system.dispatcher  
  implicit val cluster                           = Cluster(system)  
  
  AkkaManagement(system).start()  
  ClusterBootstrap(system).start()  
  
  // ...  
  
  val routes = session  
    .build[Routes].root ~ /* ... */  
  
  val bindingFuture = Http().bindAndHandle(routes, host, port).map { serverBinding =>  
    system.log.info(s"Server online at ${serverBinding.localAddress}")  
    serverBinding  
  }  
  
  // ...  
}
```

FYI: Akka Cluster

- Cluster Specification
- Node
 - A logical member of a cluster. There could be multiple nodes on a physical machine. Defined by a hostname:port:uid tuple.
- Cluster
 - A set of nodes joined together through the membership service.
- leader
 - A single node in the cluster that acts as the leader. Managing cluster convergence and membership state transitions.
- Seed Nodes
 - The seed nodes are configured contact points for new nodes joining the cluster. When a new node is started it sends a message to all seed nodes and then sends a join command to the seed node that answers first.



FYI: Akka Management

- Akka Management is a suite of tools for operating Akka Clusters.
- modules
 - akka-management: HTTP management endpoints and health checks
 - akka-managment-cluster-http: Provides HTTP endpoints for cluster monitoring and management
 - akka-managment-cluster-bootstrap: Supports cluster bootstrapping by using akka-discovery
 - akka-discovery-kubernetes-api: Module for managing k8s pod as a cluster member

Example for akka.conf(1/2)

- Example configuration for Production

```
akka {  
  cluster {  
    seed-nodes = [] # seed-nodes are empty because managed by akka-management  
    auto-down-unreachable-after = off  
  }  
  
  remote {  
    log-remote-lifecycle-events = on  
    netty.tcp {  
      hostname = "127.0.0.1"  
      hostname = ${?HOSTNAME}  
      port = 2551  
      port = ${?THREAD_WEAVER_REMOTE_PORT}  
      bind-hostname = "0.0.0.0"  
    }  
  }  
}
```

Example for akka.conf(2/2)

- Example configuration for akka-management and akka-discovery
- Configuration to find nodes from k8s pod information

```
discovery {  
  method = kubernetes-api  
  method = ${?THREAD_WEAVER_DISCOVERY_METHOD}  
  kubernetes-api {  
    pod-namespace = "thread-weaver"  
    pod-namespace = ${?THREAD_WEAVER_K8S_NAMESPACE}  
    pod-label-selector = "app=thread-weaver-api-server"  
    pod-label-selector = ${?THREAD_WEAVER_K8S_SELECTOR}  
    pod-port-name = "management"  
    pod-port-name = ${?THREAD_WEAVER_K8S_MANAGEMENT_PORT}  
  }  
}
```

```
management {  
  http {  
    hostname = "127.0.0.1"  
    hostname = ${?HOSTNAME}  
    port = 8558  
    port = ${?THREAD_WEAVER_MANAGEMENT_PORT}  
    bind-hostname = 0.0.0.0  
    bind-port = 8558  
  }  
  cluster.bootstrap {  
    contact-point-discovery {  
      discovery-method = kubernetes-api  
    }  
  }  
  contract-point {  
    fallback-port = 8558  
  }  
}
```

Deployment to EKS

FYI: Learn Kubernetes/EKS

- [Kubernetes Documentation](#)
- [Amazon EKS](#)
- [Amazon EKS Workshop](#)

build.sbt for deployment

project/plugins.sbt

```
addSbtPlugin("com.typesafe.sbt" % "sbt-native-packager" % "1.3.10")  
addSbtPlugin("com.mintbeans" % "sbt-ecr" % "0.14.1")
```

build.sbt

- Configuration for Docker

```
lazy val dockerCommonSettings = Seq(  
  dockerBaseImage := "adoptopenjdk/openjdk8:x86_64-alpine-jdk8u191-b12",  
  maintainer in Docker := "Junichi Kato <j5ik2o@gmail.com>",  
  dockerUpdateLatest := true,  
  bashScriptExtraDefines += Seq(  
    "addJava -Xms${JVM_HEAP_MIN:-1024m}",  
    "addJava -Xmx${JVM_HEAP_MAX:-1024m}",  
    "addJava -XX:MaxMetaspaceSize=${JVM_META_MAX:-512M}",  
    "addJava ${JVM_GC_OPTIONS:--XX:+UseG1GC}",  
    "addJava -Dconfig.resource=${CONFIG_RESOURCE:-application.conf}",  
    "addJava -Dakka.remote.startup-timeout=60s"  
  )  
)
```


build.sbt

- Configuration for ECR

```
val ecrSettings = Seq(  
  region in Ecr := Region.getRegion(Regions.AP_NORTHEAST_1),  
  repositoryName in Ecr := "j5ik2o/thread-weaver-api-server",  
  repositoryTags in Ecr += Seq(version.value),  
  localDockerImage in Ecr := "j5ik2o/" + (packageName in Docker).value + ":" + (version in Docker).value,  
  push in Ecr := ((push in Ecr) dependsOn (publishLocal in Docker, login in Ecr)).value  
)
```

build.sbt

```
val `api-server` = (project in file("api-server"))
  .enablePlugins(AshScriptPlugin, JavaAgent, EcrPlugin)
  .settings(baseSettings)
  .settings(dockerCommonSettings)
  .settings(ecrSettings)
  .settings(
    name := "thread-weaver-api-server",
    dockerEntrypoint := Seq("/opt/docker/bin/thread-weaver-api-server"),
    dockerUsername := Some("j5ik2o"),
    // ...
    libraryDependencies ++= Seq(
      "com.github.scopt" %% "scopt" % "4.0.0-RC2",
      "net.logstash.logback" % "logstash-logback-encoder" % "4.11" excludeAll (**/),
      "com.lightbend.akka.management" %% "akka-management" % akkaManagementVersion,
      "com.lightbend.akka.management" %% "akka-management-cluster-http" % akkaManagementVersion,
      "com.lightbend.akka.management" %% "akka-management-cluster-bootstrap" % akkaManagementVersion,
      "com.lightbend.akka.discovery" %% "akka-discovery-kubernetes-api" % akkaManagementVersion,
      "com.github.TanUkkii007" %% "akka-cluster-custom-downing" % "0.0.12",
      "com.github.everpeace" %% "healthchecks-core" % "0.4.0",
      "com.github.everpeace" %% "healthchecks-k8s-probes" % "0.4.0",
      "org.slf4j" % "jul-to-slf4j" % "1.7.26",
      "ch.qos.logback" % "logback-classic" % "1.2.3",
      "org.codehaus.janino" % "janino" % "3.0.6"
    )
  )
```

Deployment to Local Cluster(minikube)

Deployment to Local Cluster(minikube)

- start minikube
- Helm Implementation
- Create namespaces and service accounts
- Deploy DB
- Create Schema
- Building an image for your application
- Deploy the image for the application

```
$ minikube start --vmdriver virtualbox \
  --kubernetes-version v1.12.8 --cpus 6 --memory 5000 --disk-size 30g
$ helm init
$ kubectl create namespace thread-weaver
$ kubectl create serviceaccount thread-weaver
$ helm install ./mysql --namespace thread-weaver \
  -f ./mysql/environments/${ENV_NAME}-values.yaml
$ helm install ./dynamodb --namespace thread-weaver \
  -f ./dynamodb/environments/${ENV_NAME}-values.yaml
$ sbt -Dmysql.host="$(minikube ip)" -Dmysql.port=30306 \
  'migrate-mysql/run'
$ DYNAMODB_HOST="$(minikube ip)" DYNAMODB_PORT=32000 \
  sbt 'migrate-dynamodb/run'
$ eval $(minikube docker-env)
$ sbt api-server/docker:publishLocal
$ helm install ./thread-weaver-api-server \
  --namespace thread-weaver \
  -f ./thread-weaver-api-server/environments/${ENV_NAME}-values.yaml
```

Helm charts

FYI: Helm

- Helm
 - Helm is 'The package manager for Kubernetes'
- Charts
 - Helm uses a packaging format called charts
- CLI
 - `helm init`
 - initialize Helm on both client and server(tiller)
 - `helm package`
 - package a chart directory into a chart archive
 - `helm install`
 - install a chart archive
 - `helm upgrade`
 - upgrade a release
 - `helm rollback`
 - roll back a release to a previous revision



deployment.yaml(1/2)

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: {{ template "name" . }}
spec:
  replicas: {{ .Values.replicaCount }}
  selector:
    matchLabels:
      app: {{ template "name" . }}
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 0
  template:
    metadata:
      labels:
        app: {{ template "name" . }}
    spec:
      containers:
        - image: "{{ .Values.image.repository }}:{{ .Values.image.tag }}"
          imagePullPolicy: {{ .Values.image.pullPolicy }}
          name: {{ template "name" . }}
```

```
env:
  - name: AWS_REGION
    value: "ap-northeast-1"
  - name: HOSTNAME
    valueFrom:
      fieldRef:
        apiVersion: v1
        fieldPath: status.podIP
  - name: ENV_NAME
    value: {{ .Values.envName | quote }}
  - name: CONFIG_RESOURCE
    value: {{ .Values.configResource | quote }}
  - name: JVM_HEAP_MIN
    value: {{ .Values.jvmHeapMin | quote }}
  - name: JVM_HEAP_MAX
    value: {{ .Values.jvmHeapMax | quote }}
  - name: JVM_META_MAX
    value: {{ .Values.jvmMetaMax | quote }}
```

deployment.yaml(2/2)

```

- name: THREAD_WEAVER_SLICK_URL
  value: {{.Values.db.url | quote}}
- name: THREAD_WEAVER_SLICK_USER
  value: {{.Values.db.user | quote}}
- name: THREAD_WEAVER_SLICK_PASSWORD
  valueFrom:
    secretKeyRef:
      name: thread-weaver-app-secrets
      key: mysql.password
- name: THREAD_WEAVER_SLICK_MAX_POOL_SIZE
  value: {{.Values.db.maxPoolSize | quote}}
- name: THREAD_WEAVER_SLICK_MIN_IDLE_SIZE
  value: {{.Values.db.minIdleSize | quote}}
ports:
- name: remoting
  containerPort: 2551
- name: {{.Values.service.name }}
  containerPort: {{.Values.service.internalPort }}
- name: management
  containerPort: 8558

```

```

readinessProbe:
  tcpSocket:
    port: 18080
    initialDelaySeconds: 60
    periodSeconds: 30
livenessProbe:
  tcpSocket:
    port: 18080
    initialDelaySeconds: 60
    periodSeconds: 30

```

- Describes the settings for starting the container, such as image name, tags, environment variables, and ports

service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: {{ template "name" . }}
  labels:
    app: {{ template "name" . }}
    chart: {{ .Chart.Name }}-{{ .Chart.Version | replace "+" "_" }}
    release: {{ .Release.Name }}
    heritage: {{ .Release.Service }}
spec:
  selector:
    app: {{ template "name" . }}
  type: {{ .Values.service.type }}
  ports:
    - protocol: TCP
      name: api
      port: 8080
      targetPort: api
    - protocol: TCP
      name: management
      port: 8558
      targetPort: management
```

- Build a Service of type LoadBalancer to make it externally accessible

rbac.yaml

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: thread-weaver-api-server
rules:
- apiGroups: ["" ]
  resources: ["pods"]
  verbs: ["get", "watch", "list"]
---
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: thread-weaver-api-server
subjects:
- kind: User
  name: system:serviceaccount:thread-weaver:default
roleRef:
  kind: Role
  name: thread-weaver-api-server
  apiGroup: rbac.authorization.k8s.io
```

Configure RBAC so that akka-discovery can find the nodes

Verification for Local Cluster

```
API_HOST=$(minikube ip)
API_PORT=$(kubectl get svc thread-weaver-api-server -n thread-weaver -ojsonpath="{.spec.ports[?(@.name==\"api\")].port}")

ACCOUNT_ID=01DB5QXD4NP0XQTV92K42B3XBF
ADMINISTRATOR_ID=01DB5QXD4NP0XQTV92K42B3XBF

THREAD_ID=$(curl -v -X POST "http://$API_HOST:$API_PORT/v1/threads/create" -H "accept: application/json" -H "Content-Type: application/json" \
  -d "{\"accountId\":\"${ACCOUNT_ID}\",\"title\":\"string\",\"remarks\":\"string\",\"administratorIds\":[${ADMINISTRATOR_ID}],\""}" \
  && echo "THREAD_ID=$THREAD_ID")
sleep 3
curl -v -X GET "http://$API_HOST:$API_PORT/v1/threads/${THREAD_ID}?account_id=${ACCOUNT_ID}" -H "accept: application/json"
```

Deployment to Production Cluster(EKS)

Build Kubernetes Cluster(EKS)

- Build the required components for the EKS cluster in advance
 - subnet
 - security group
 - internet-gw
 - eip
 - nat-gw
 - route table
 - ecr
- Build the database required by the application
 - rds(aurora)
 - dynamodb(with shema)

```
$ terraform plan  
$ terraform apply
```

Build Kubernetes Cluster(EKS)

- Build an EKS cluster
 - eksctl

```
$ eksctl create cluster \  
  --name ${CLUSTER_NAME} \  
  --region ${AWS_REGION} \  
  --nodes ${NODES} \  
  --nodes-min ${NODES_MIN} \  
  --nodes-max ${NODES_MAX} \  
  --node-type ${INSTANCE_TYPE} \  
  --full-ecr-access \  
  --node-ami ${NODE_AMI} \  
  --version ${K8S_VERSION} \  
  --nodegroup-name ${NODE_GROUP_NAME} \  
  --vpc-private-subnets=${SUBNET_PRIVATE1},${SUBNET_PRIVATE2},${SUBNET_PRIVATE3} \  
  --vpc-public-subnets=${SUBNET_PUBLIC1},${SUBNET_PUBLIC2},${SUBNET_PUBLIC3}
```

- Initial Setup (RBAC settings, etc.)

```
tools/eks/helm $ kubectl apply -f ./rbac-config.yaml  
$ helm init  
$ kubectl create namespace thread-weaver  
$ kubectl create serviceaccount thread-weaver  
tools/deploy/eks $ kubectl apply -f secret.yaml
```

Deploy to Kubernetes Cluster(EKS)

- docker build & push to ecr

```
$ AWS_DEFAULT_PROFILE=xxxxx sbt api-server/ecr:push
```

- flyway migrate(should be implemented as k8s job)

```
$ docker run --rm -v $(pwd)/tools/flyway/src/test/resources/db-migration:/flyway/sql -v $(pwd):/flyway/conf boxfuse/flyway migrate
```

- deploy

```
$ helm install ./thread-weaver-api-server \
--namespace thread-weaver \
-f ./thread-weaver-api-server/environments/${ENV_NAME}-values.yaml
```

Verification for Kubernetes Cluster(EKS)

```
API_HOST=$(kubectl get svc thread-weaver-api-server -n thread-weaver -ojsonpath="{.status.loadBalancer.ingress[0].hostname}")
API_PORT=$(kubectl get svc thread-weaver-api-server -n thread-weaver -ojsonpath="{.spec.ports[?(@.name==\"api\")].port}")

ACCOUNT_ID=01DB5QXD4NP0XQTV92K42B3XBF
ADMINISTRATOR_ID=01DB5QXD4NP0XQTV92K42B3XBF

THREAD_ID=$(curl -v -X POST "http://$API_HOST:$API_PORT/v1/threads/create"
-H "accept: application/json" -H "Content-Type: application/json" \
-d "{\"accountId\":\"${ACCOUNT_ID}\", ... \"memberIds\":[${ACCOUNT_ID}],\"createAt\":10000}" | jq -r .threadId)
echo "THREAD_ID=$THREAD_ID"
sleep 3
curl -v -X GET "http://$API_HOST:$API_PORT/v1/threads/${THREAD_ID}?account_id=${ACCOUNT_ID}" -H "accept: application/json"
```


Perspective to be considered for production operations

- Event Schema Evolution
 - [Persistence - Schema Evolution](#)
- Split-Brain Resolver
 - [Split Brain Resolver](#)
 - [TanUkkii007/akka-cluster-custom-downing](#)
 - [Akka Clusterの耐障害設計](#)
- Distributed Tracing
 - [kamon-io/Kamon](#)
 - [alevkhomich/akka-tracing](#)

Summary

- ReadModelUpdater needs to be changed in design so that it can be operated on another node
- Use multi-jvm testing or pods on minikube to validate sharded actors
- Split Brain Resolver is available in OSS as well as commercial versions
- Pods on EKS must be managed as cluster members. Issues remain with cluster new version migration

We are hiring engineers who work together
<http://corp.chatwork.com/ja/recruit/>

